

Hierarchical Leader Election Algorithm With Remoteness Constraint

Mohamed Tbarka

November 11, 2019



Outline	Introduction	Preliminaries	H. Leader Election Algorithm	Correctness	Implementation	Conclusion
●	○ ○ ○ ○	○ ○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○ ○ ○	○	○ ○ ○ ○ ○	○ ○

Outline

1 Introduction

- Leader Election Problem
- State of art

2 Preliminaries

- System Model
- Problem Definition

3 H. Leader Election Algorithm

- Informal Description
- Nodes, Neighbors and Heights
- Initial State
- Description Of The Algorithm
- Sample Execution

4 Correctness

Outline

- 1 Introduction
 - Leader Election Problem
 - State of art
- 2 Preliminaries
 - System Model
 - Problem Definition
- 3 H. Leader Election Algorithm
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Description Of The Algorithm
 - Sample Execution
- 4 Correctness

What is the problem of leader election ?

Leader election is an important primitive for distributed computing, useful as a sub-routine for any application that requires the selection of a unique processor among multiple candidate processors (video conferencing, multi-player games, ...).

What is distributed computing ?

Distributed computing is a model in which components of a software system are shared among multiple computers to improve efficiency and performance.

State of art




Figure:

Outline

- 1 Introduction
 - Leader Election Problem
 - State of art
- 2 Preliminaries
 - System Model
 - Problem Definition
- 3 H. Leader Election Algorithm
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Description Of The Algorithm
 - Sample Execution
- 4 Correctness

System Model

- The system is consisting of a set P of computing nodes and a set χ of directed communication channels from one node to another node.

System Model

- The system is consisting of a set P of computing nodes and a set χ of directed communication channels from one node to another node.
- The whole system as a set of (infinite) state machines that interact through shared events.

System Model

- The system is consisting of a set P of computing nodes and a set χ of directed communication channels from one node to another node.

System Model

- The system is consisting of a set P of computing nodes and a set χ of directed communication channels from one node to another node.
- The whole system as a set of (infinite) state machines that interact through shared events.

Asynchronous Dynamic Links' Model

The state of $Channel(u, v)$, which models the communication channel from node u to node v , consists of:

- a $status_{uv}$ variable;

Asynchronous Dynamic Links' Model

The state of $Channel(u, v)$, which models the communication channel from node u to node v , consists of:

- a $status_{uv}$ variable;
- and a queue $mqueue_{uv}$ of messages.

Configurations & Executions

- The notion of configuration is used to capture an instantaneous snapshot of the state of the entire system.

Configurations & Executions

- The notion of configuration is used to capture an instantaneous snapshot of the state of the entire system.
- A configuration is a vector of node states, one for each node in P , and a vector of channel states, one for each channel in χ .

Configurations & Executions

In an initial configuration:

- each node is in an initial state (according to its algorithm),

Configurations & Executions

In an initial configuration:

- each node is in an initial state (according to its algorithm),
- for each channel $Channel(u, v)$, $mqueue_{uv}$ is empty, and

Configurations & Executions

In an initial configuration:

- each node is in an initial state (according to its algorithm),
- for each channel $Channel(u, v)$, $mqueue_{uv}$ is empty, and
- for all nodes u and v , $status_{uv} = status_{vu}$ (i.e., either both channels between u and v are up, or both are down).

Configurations & Executions

An execution is an infinite sequence $C_0, e_1, C_1, e_2, C_2, \dots$ of alternating configurations and events, starting with an initial configuration and, if finite, ending with a configuration.

Problem Definition

Each node u in the system has :

- a local variable lid_u to hold the id of the supreme leader;

Problem Definition

Each node u in the system has :

- a local variable lid_u to hold the id of the supreme leader;
- another local variable $slid_u$ to hold the identifier of the sub-leader whose remoteness towards u obeys the constraint.

Outline

- 1 Introduction
 - Leader Election Problem
 - State of art
- 2 Preliminaries
 - System Model
 - Problem Definition
- 3 H. Leader Election Algorithm**
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Description Of The Algorithm
 - Sample Execution
- 4 Correctness

Informal description

After a leader is gone, the algorithm consists on three waves:

- First wave : initiated by one of the lost leader's neighbors looking for it;

Informal description

After a leader is gone, the algorithm consists on three waves:

- First wave : initiated by one of the lost leader's neighbors looking for it;
- Second wave : initiated by the node located at the edge of the network if the search has hit a dead-end;

Informal description

After a leader is gone, the algorithm consists on three waves:

- First wave : initiated by one of the lost leader's neighbors looking for it;
- Second wave : initiated by the node located at the edge of the network if the search has hit a dead-end;
- Third wave : initiated by the same node which initiated the first wave updating the other nodes' heights and constructing the spanning tree.

Nodes, Neighbors and Heights

- When a node u gets a *ChannelUp* event for the channel from u to v , it puts v in a local set variable called *forming_u*.

And when u gets a *ChannelDown* event for the channel from u to v , it

Nodes, Neighbors and Heights

- When a node u gets a *ChannelUp* event for the channel from u to v , it puts v in a local set variable called $forming_u$.
- When u subsequently receives a message from v , it moves v from its $forming_u$ set to a local set variable called N_u (N for neighbor).

And when u gets a *ChannelDown* event for the channel from u to v , it

Nodes, Neighbors and Heights

- When a node u gets a *ChannelUp* event for the channel from u to v , it puts v in a local set variable called $forming_u$.
- When u subsequently receives a message from v , it moves v from its $forming_u$ set to a local set variable called N_u (N for neighbor).
- If u gets a message from a node which is neither in its forming set, nor in N_u , it ignores that message.

And when u gets a *ChannelDown* event for the channel from u to v , it

Nodes, Neighbors and Heights

- When a node u gets a *ChannelUp* event for the channel from u to v , it puts v in a local set variable called $forming_u$.
- When u subsequently receives a message from v , it moves v from its $forming_u$ set to a local set variable called N_u (N for neighbor).
- If u gets a message from a node which is neither in its forming set, nor in N_u , it ignores that message.
- And when u gets a *ChannelDown* event for the channel from u to v , it removes v from $forming_u$ or N_u , as appropriate.

And when u gets a *ChannelDown* event for the channel from u to v , it

Initial State

- $forming_u$ is empty,
 N_u equals the set of neighbors of u in G_{chan}^{init}

Initial State

- $forming_u$ is empty,
 N_u equals the set of neighbors of u in G_{chan}^{init}
- $height_u[u] = (0, 0, 0, \delta_u, 0, l, u)$ where l is the *id* of a fixed node in u 's connected component in G_{chan}^{init} (the current leader), and δ_u equals the distance from u to l in G_{chan}^{init} ,

Initial State

- $forming_u$ is empty,
 N_u equals the set of neighbors of u in G_{chan}^{init}
- $height_u[u] = (0, 0, 0, \delta_u, 0, l, u)$ where l is the *id* of a fixed node in u 's connected component in G_{chan}^{init} (the current leader), and δ_u equals the distance from u to l in G_{chan}^{init} ,
- for each v in N_u , $height_u[v] = height_v[v]$ (i.e., u has accurate information about v 's height), and

Initial State

- $forming_u$ is empty,
 N_u equals the set of neighbors of u in G_{chan}^{init}
- $height_u[u] = (0, 0, 0, \delta_u, 0, l, u)$ where l is the *id* of a fixed node in u 's connected component in G_{chan}^{init} (the current leader), and δ_u equals the distance from u to l in G_{chan}^{init} ,
- for each v in N_u , $height_u[v] = height_v[v]$ (i.e., u has accurate information about v 's height), and
- \mathcal{T}_u is initialized properly with respect to the definition of causal clocks.

Heights

The height for each node is a 7-tuple of integers $((\tau, oid, r), \delta, (nlts, lid), id)$, where the first three components are referred to as the reference level (RL) and the fifth and sixth.

- τ , a non-negative timestamp which is either 0 or the value of the causal clock time when the current search for an alternate path to the leader was initiated.

Heights

The height for each node is a 7-tuple of integers $((\tau, oid, r), \delta, (nlts, lid), id)$, where the first three components are referred to as the reference level (RL) and the fifth and sixth.

- τ , a non-negative timestamp which is either 0 or the value of the causal clock time when the current search for an alternate path to the leader was initiated.
- oid , is a non-negative value that is either 0 or the id of the node that started the current search (we assume node ids are positive integers).

Heights

The height for each node is a 7-tuple of integers $((\tau, oid, r), \delta, (nlts, lid), id)$, where the first three components are referred to as the reference level (RL) and the fifth and sixth.

- τ , a non-negative timestamp which is either 0 or the value of the causal clock time when the current search for an alternate path to the leader was initiated.
- oid , is a non-negative value that is either 0 or the id of the node that started the current search (we assume node ids are positive integers).
- r , a bit that is set to 0 when the current search is initiated and set to 1 when the current search hits a dead end.

Heights

The height for each node is a 7-tuple of integers $((\tau, oid, r), \delta, (nlts, lid), id)$, where the first three components are referred to as the reference level (RL) and the fifth and sixth.

- τ , a non-negative timestamp which is either 0 or the value of the causal clock time when the current search for an alternate path to the leader was initiated.
- oid , is a non-negative value that is either 0 or the id of the node that started the current search (we assume node ids are positive integers).
- r , a bit that is set to 0 when the current search is initiated and set to 1 when the current search hits a dead end.
- δ

Subroutines

ELECTSELF

1. $height[u] := (0, 0, 0, 0, -\mathcal{T}_u, u, u)$

REFLECTREFLEVEL

1. $height[u] := (\tau, oid, 1, 0, nlts^u, lid^u, u)$

PROPAGATELARGESTREFLEVEL

1. $(\tau^u, oid^u, r^u) := \max\{(\tau^w, oid^w, r^w) \mid w \in N\}$
2. $\delta^u := \min\{ \delta^w \mid w \in N \text{ and } (\tau^u, oid^u, r^u) = (\tau^w, oid^w, r^w) \} - 1$

STARTNEWREFLEVEL

1. $height[u] := (\mathcal{T}_u, u, 0, 0, nlts^u, lid^u, u)$

ADOPTLPIFPRIORITY(v)

1. if $((nlts^v < nlts^u) \text{ or } ((nlts^v = nlts^u) \text{ and } (lid^v < lid^u)))$
2. $height[u] := (\tau^v, oid^v, r^v, \delta^v + 1, nlts^v, lid^v, u)$
3. else send Update($height[u]$) to v
4. end if

Outline

- 1 Introduction
 - Leader Election Problem
 - State of art
- 2 Preliminaries
 - System Model
 - Problem Definition
- 3 H. Leader Election Algorithm
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Description Of The Algorithm
 - Sample Execution
- 4 **Correctness**

Outline

- 1 Introduction
 - Leader Election Problem
 - State of art
- 2 Preliminaries
 - System Model
 - Problem Definition
- 3 H. Leader Election Algorithm
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Description Of The Algorithm
 - Sample Execution
- 4 Correctness


What's JBotSim ?

JBoTSim is a java library that offers basic primitives for proto-typing, running, and visualizing distributed algorithms in dynamic networks.


Outline	Introduction	Preliminaries	H. Leader Election Algorithm	Correctness	Implementation	Conclusion
○	○ ○ ○	○ ○ ○ ○ ○ ○ ○	○ ○ ○ ○ ○ ○ ○ ○ ○ ○ ○	○	○ ○ ● ○	○ ○ ○

Simulation

How is our algorithm's performance ?



(a) Latency



(b) Resilience

Figure: Simulation Results

Outline

- 1 Introduction
 - Leader Election Problem
 - State of art
- 2 Preliminaries
 - System Model
 - Problem Definition
- 3 H. Leader Election Algorithm
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Description Of The Algorithm
 - Sample Execution
- 4 Correctness

Is The Algorithm Perfect ?

An open question is how to extend our algorithm and its analysis to handle a wider range of clocks, such as approximately synchronized clocks and vector clocks.

Is The Algorithm Perfect ?

Question ?