

Hierarchical Leader Election Algorithm With Remoteness Constraint

Mohamed Tbarka

November 15, 2019



Outline	Introduction	Preliminaries	H. Leader Election Algorithm	Correctness	Implementation	Conclusion
●	○	○	○	○○	○	○
	○○	○○○○	○○○	○○	○	○○
	○	○	○○○○	○○	○○○○	
			○○○○	○○		
			○○○○	○○○○		
			○○○○			

Outline

1 Introduction

- Leader Election Problem
- State of art

2 Preliminaries

- System Model
- Problem Definition

3 H. Leader Election Algorithm

- Informal Description
- Nodes, Neighbors and Heights
- Initial State
- Algorithm Description

4 Correctness

- Bounding the Number of Elections
- Bounding the Number of New Reference Levels
- Bounding the Number of Messages
- Proving The Leader-Oriented DAG Aspect

5 Implementation

- The Tool Used
- Simulation
- Performance Test

6 Conclusion

- Is The Algorithm Perfect ?

Outline	Introduction	Preliminaries	H. Leader Election Algorithm	Correctness	Implementation	Conclusion
o	●	o oo o	o oooo o oooo oooo	oo o o o	o o o oooo	o oo
o	●	ooooo	oooo	o	o	o
o	●	o	o	o	o	o
o	●	o	oooo	o	o	oo

Outline

1 Introduction

- Leader Election Problem
- State of art

2 Preliminaries

- System Model
- Problem Definition

3 H. Leader Election Algorithm

- Informal Description
- Nodes, Neighbors and Heights
- Initial State
- Algorithm Description

4 Correctness

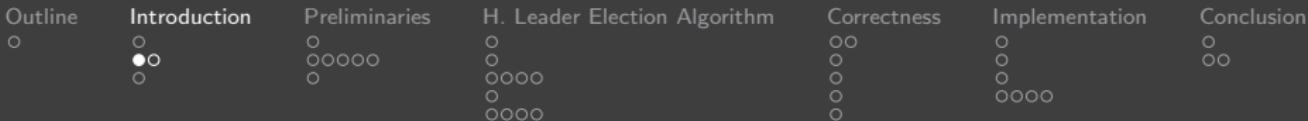
- Bounding the Number of Elections
- Bounding the Number of New Reference Levels
- Bounding the Number of Messages
- Proving The Leader-Oriented DAG Aspect

5 Implementation

- The Tool Used
- Simulation
- Performance Test

6 Conclusion

- Is The Algorithm Perfect ?

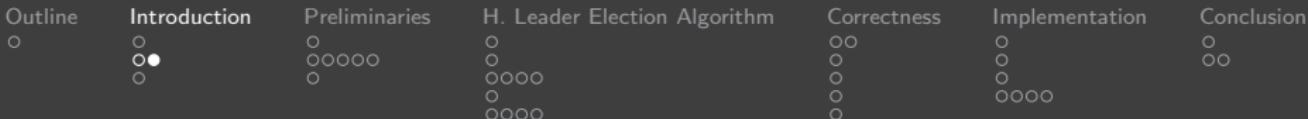


Leader Election Problem

What is the problem of leader election ?

Leader election is an important primitive for distributed computing, useful as a sub-routine for any application that requires the selection of a unique processor among multiple candidate processors (video conferencing, multi-player games, ...).



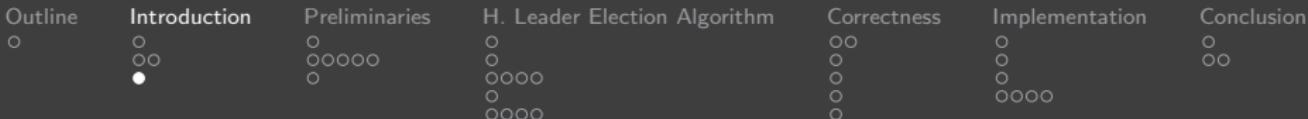


Leader Election Problem

What is distributed computing ?

Distributed computing is a model in which components of a software system are shared among multiple computers to improve efficiency and performance.





State of art

State of art

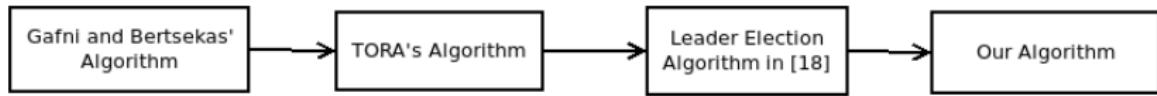


Figure: State Of Art

Outline	Introduction	Preliminaries	H. Leader Election Algorithm	Correctness	Implementation	Conclusion
o	o	●	o	oo	o	o
o	oo	oooo	o	o	o	oo
o	o	o	oooo	o	o	
			oooo	o	oooo	

Outline

1 Introduction

- Leader Election Problem
- State of art

2 Preliminaries

- System Model
- Problem Definition

3 H. Leader Election Algorithm

- Informal Description
- Nodes, Neighbors and Heights
- Initial State
- Algorithm Description

4 Correctness

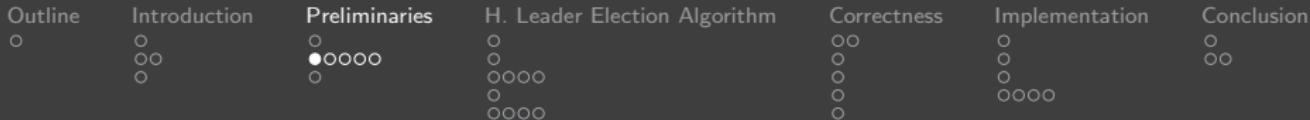
- Bounding the Number of Elections
- Bounding the Number of New Reference Levels
- Bounding the Number of Messages
- Proving The Leader-Oriented DAG Aspect

5 Implementation

- The Tool Used
- Simulation
- Performance Test

6 Conclusion

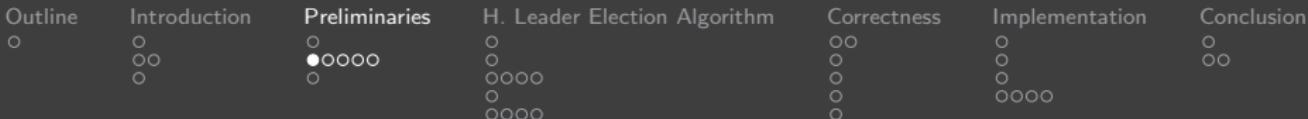
- Is The Algorithm Perfect ?



System Model

- The system is consisting of a set P of computing nodes and a set χ of directed communication channels from one node to another node.

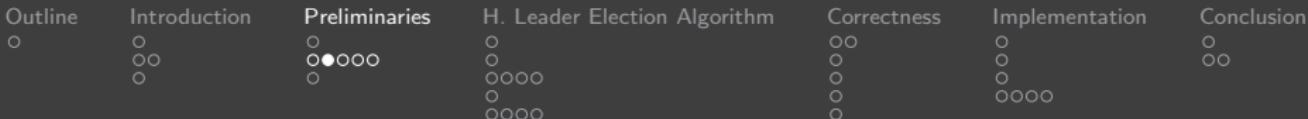




System Model

- The system is consisting of a set P of computing nodes and a set χ of directed communication channels from one node to another node.
- The whole system as a set of (infinite) state machines that interact through shared events.





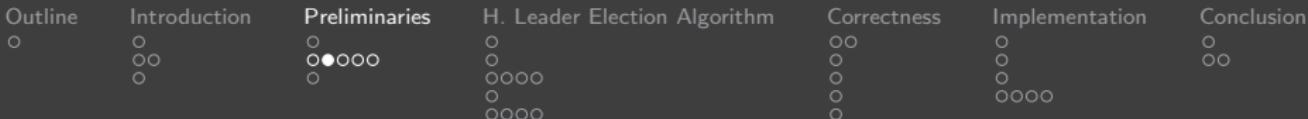
System Model

Asynchronous Dynamic Links' Model

The state of $\text{Channel}(u, v)$, which models the communication channel from node u to node v , consists of:

- a status_{uv} variable;





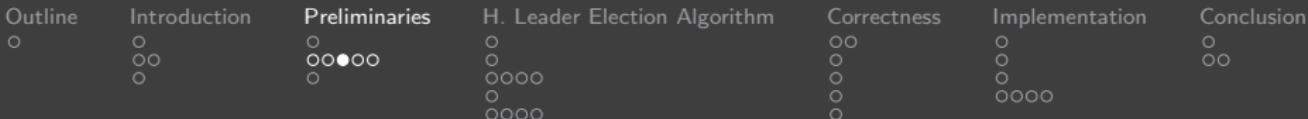
System Model

Asynchronous Dynamic Links' Model

The state of $\text{Channel}(u, v)$, which models the communication channel from node u to node v , consists of:

- a status_{uv} variable;
- and a queue $mqueue_{uv}$ of messages.



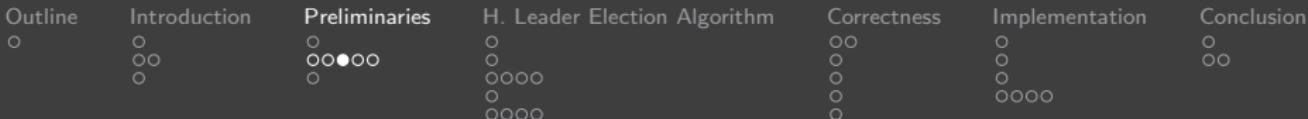


System Model

Configurations & Executions

- The notion of configuration is used to capture an instantaneous snapshot of the state of the entire system.



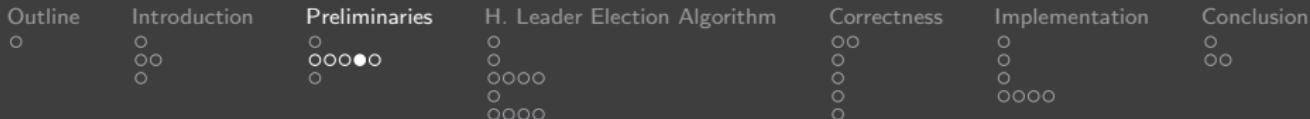


System Model

Configurations & Executions

- The notion of configuration is used to capture an instantaneous snapshot of the state of the entire system.
- A configuration is a vector of node states, one for each node in P , and a vector of channel states, one for each channel in χ .





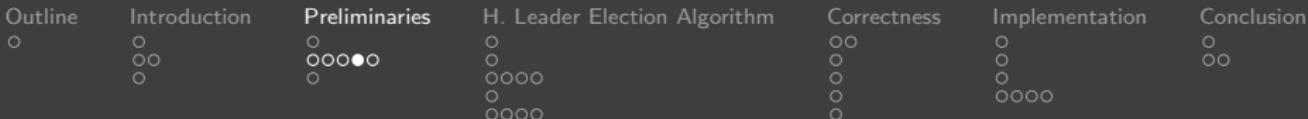
System Model

Configurations & Executions

In an initial configuration:

- each node is in an initial state (according to its algorithm),





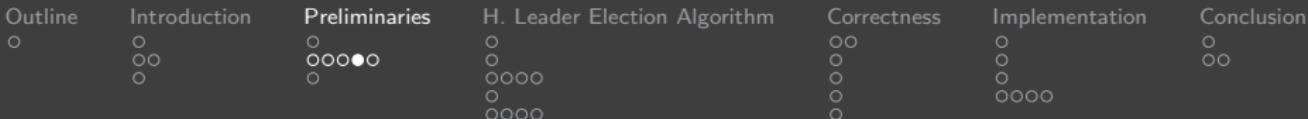
System Model

Configurations & Executions

In an initial configuration:

- each node is in an initial state (according to its algorithm),
- for each channel $Channel(u, v)$, $mqueue_{uv}$ is empty, and





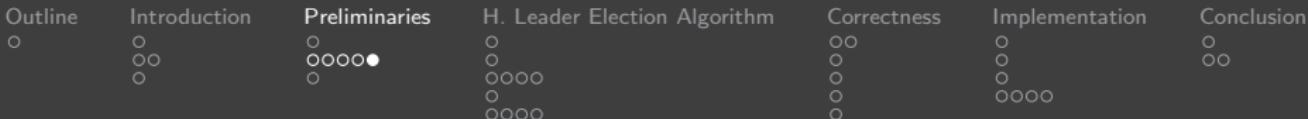
System Model

Configurations & Executions

In an initial configuration:

- each node is in an initial state (according to its algorithm),
- for each channel $Channel(u, v)$, $mqueue_{uv}$ is empty, and
- for all nodes u and v , $status_{uv} = status_{vu}$ (i.e., either both channels between u and v are up, or both are down).



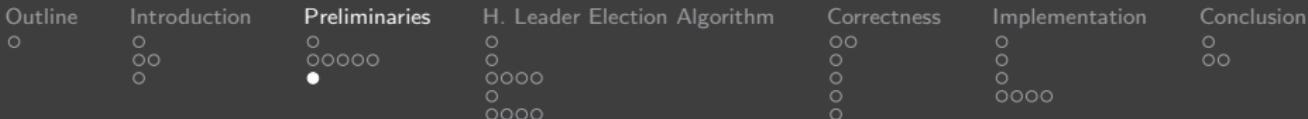


System Model

Configurations & Executions

An execution is an infinite sequence $C_0, e_1, C_1, e_2, C_2, \dots$ of alternating configurations and events, starting with an initial configuration and, if finite, ending with a configuration.





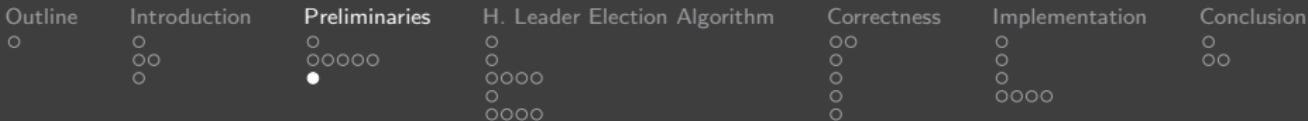
Problem Definition

Problem Definition

Each node u in the system should have after the last topology change :

- a local variable lid_u to hold the id of the supreme leader;





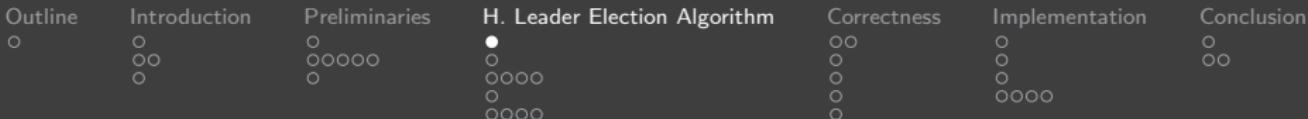
Problem Definition

Problem Definition

Each node u in the system should have after the last topology change :

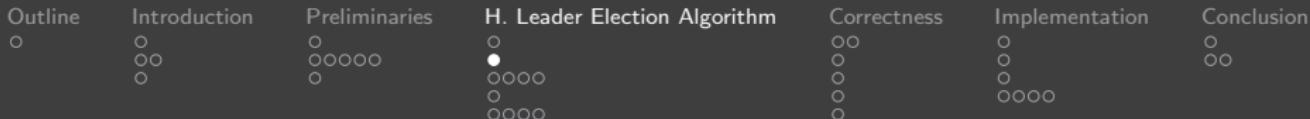
- a local variable lid_u to hold the id of the supreme leader;
- another local variable $slid_u$ to hold the identifier of the sub-leader whose remoteness towards u obeys the constraint.





Outline

- 1** Introduction
 - Leader Election Problem
 - State of art
- 2** Preliminaries
 - System Model
 - Problem Definition
- 3** H. Leader Election Algorithm
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Algorithm Description
- 4** Correctness
 - Bounding the Number of Elections
 - Bounding the Number of New Reference Levels
 - Bounding the Number of Messages
 - Proving The Leader-Oriented DAG Aspect
- 5** Implementation
 - The Tool Used
 - Simulation
 - Performance Test
- 6** Conclusion
 - Is The Algorithm Perfect ?

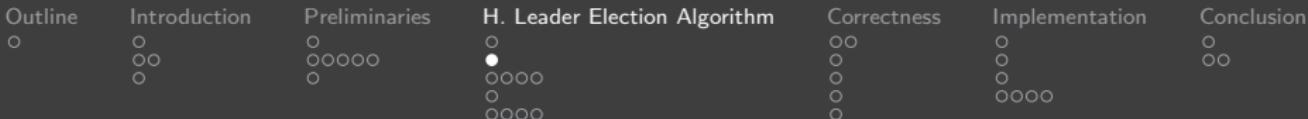


Informal Description

Informal description

After a leader is gone, the algorithm consists on three waves:





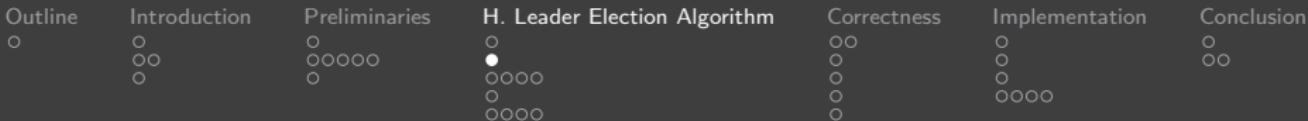
Informal Description

Informal description

After a leader is gone, the algorithm consists on three waves:

- First wave : initiated by one of the lost leader's neighbors looking for it;





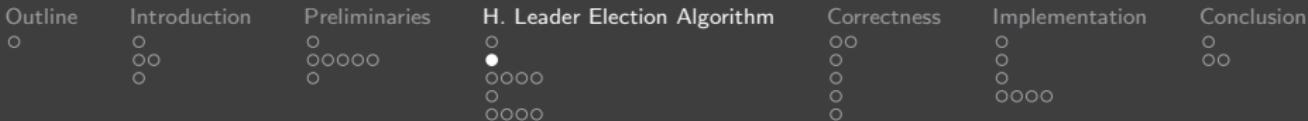
Informal Description

Informal description

After a leader is gone, the algorithm consists on three waves:

- First wave : initiated by one of the lost leader's neighbors looking for it;
- Second wave : initiated by the node located at the edge of the network if the search has hit a dead-end;





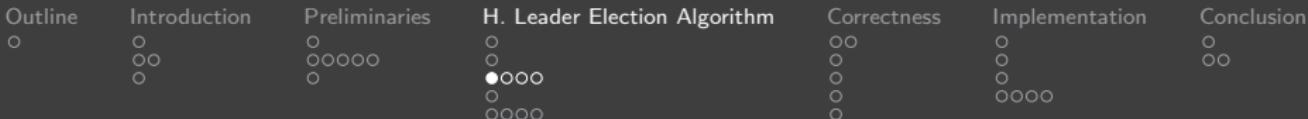
Informal Description

Informal description

After a leader is gone, the algorithm consists on three waves:

- First wave : initiated by one of the lost leader's neighbors looking for it;
- Second wave : initiated by the node located at the edge of the network if the search has hit a dead-end;
- Third wave : initiated by the same node which initiated the first wave updating the other nodes' heights and constructing the spanning tree.



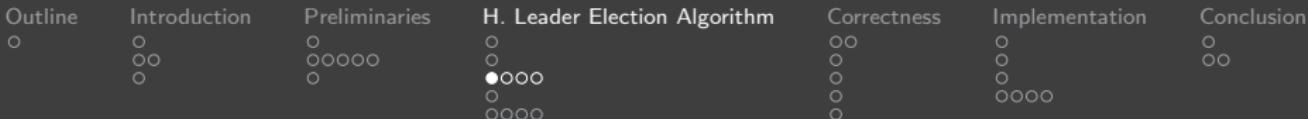


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- When a node u gets a *ChannelUp* event for the channel from u to v , it puts v in a local set variable called $forming_u$.



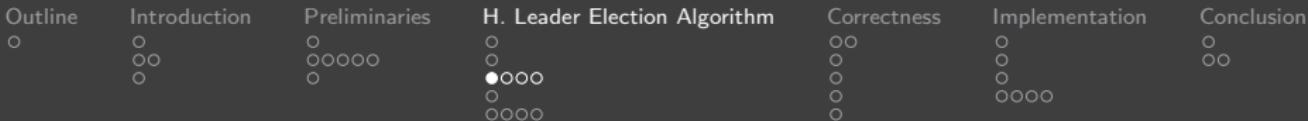


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- When a node u gets a *ChannelUp* event for the channel from u to v , it puts v in a local set variable called $forming_u$.
- When u subsequently receives a message from v , it moves v from its $forming_u$ set to a local set variable called N_u (N for neighbor).



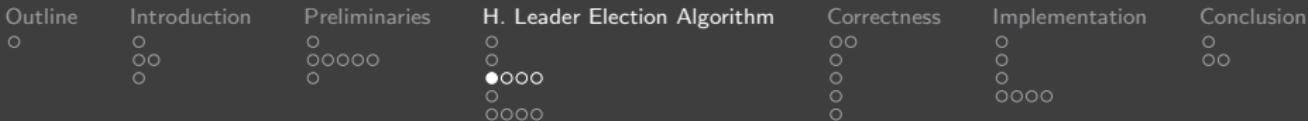


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- When a node u gets a *ChannelUp* event for the channel from u to v , it puts v in a local set variable called $forming_u$.
- When u subsequently receives a message from v , it moves v from its $forming_u$ set to a local set variable called N_u (N for neighbor).
- If u gets a message from a node which is neither in its forming set, nor in N_u , it ignores that message.



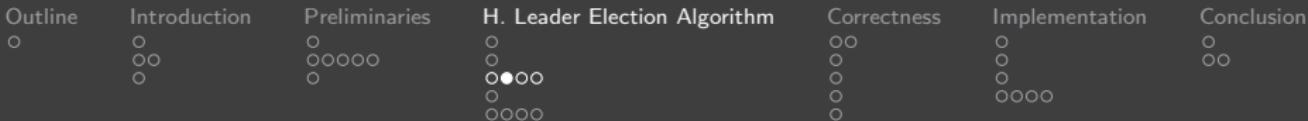


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- When a node u gets a *ChannelUp* event for the channel from u to v , it puts v in a local set variable called $forming_u$.
- When u subsequently receives a message from v , it moves v from its $forming_u$ set to a local set variable called N_u (N for neighbor).
- If u gets a message from a node which is neither in its forming set, nor in N_u , it ignores that message.
- And when u gets a *ChannelDown* event for the channel from u to v , it removes v from $forming_u$ or N_u , as appropriate.



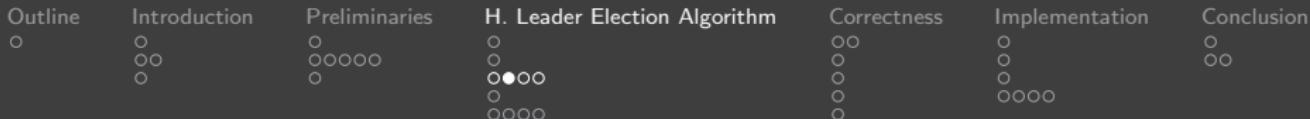


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- Nodes assign virtual directions to their links using variables called heights.

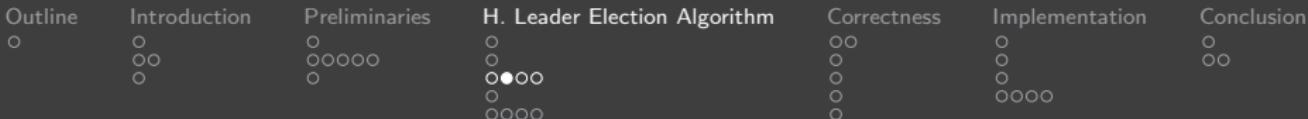




Nodes, Neighbors and Heights

- Nodes assign virtual directions to their links using variables called heights.
- For each link (u, v) of node u , u considers the link as incoming if the height that u has recorded for v is larger than u 's own height; otherwise u considers the link as outgoing (directed from u to v).



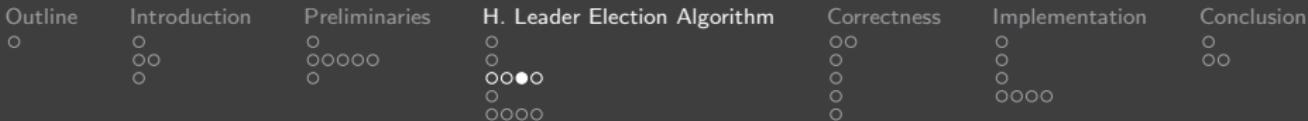


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- Nodes assign virtual directions to their links using variables called heights.
- For each link (u, v) of node u , u considers the link as incoming if the height that u has recorded for v is larger than u 's own height; otherwise u considers the link as outgoing (directed from u to v).
- Heights are compared using lexicographic ordering.





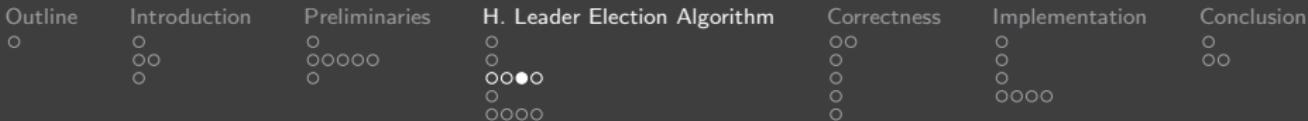
Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

The height for each node is a 7-tuple of integers
 $((\tau, oid, r), \delta, (nlts, lid), id)$ where :

- τ , a non-negative timestamp which is either 0 or the value of the causal clock time when the current search for an alternate path to the leader was initiated.





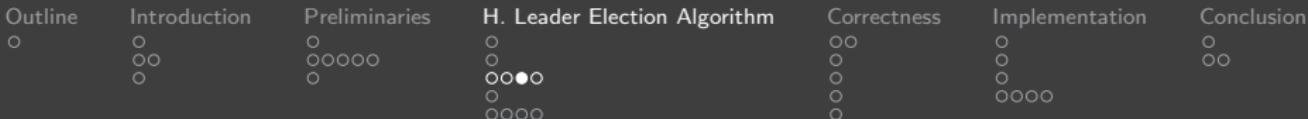
Nodes, Neighbors and Heights

The height for each node is a 7-tuple of integers

$((\tau, oid, r), \delta, (nlts, lid), id)$ where :

- τ , a non-negative timestamp which is either 0 or the value of the causal clock time when the current search for an alternate path to the leader was initiated.
- oid , is a non-negative value that is either 0 or the id of the node that started the current search (we assume node ids are positive integers).





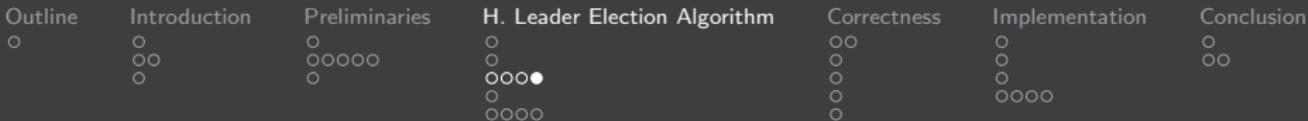
Nodes, Neighbors and Heights

The height for each node is a 7-tuple of integers

$((\tau, oid, r), \delta, (nlts, lid), id)$ where :

- τ , a non-negative timestamp which is either 0 or the value of the causal clock time when the current search for an alternate path to the leader was initiated.
- oid , is a non-negative value that is either 0 or the id of the node that started the current search (we assume node ids are positive integers).
- r , a bit that is set to 0 when the current search is initiated and set to 1 when the current search hits a dead end.



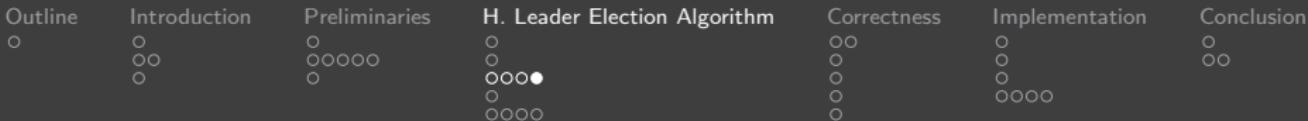


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- δ , an integer that is set to ensure that links are directed appropriately to neighbors with the same first three components.



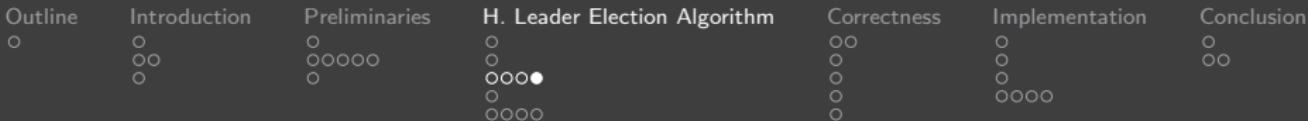


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- δ , an integer that is set to ensure that links are directed appropriately to neighbors with the same first three components.
- nlt , a non-positive timestamp whose absolute value is the causal clock time when the current leader was elected.



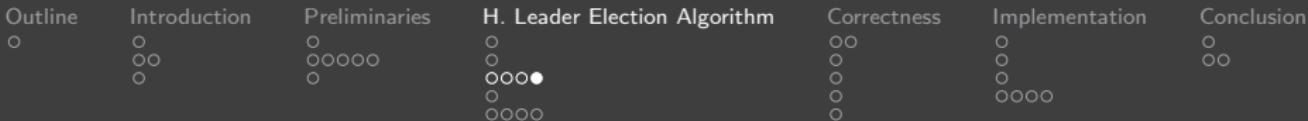


Nodes, Neighbors and Heights

Nodes, Neighbors and Heights

- δ , an integer that is set to ensure that links are directed appropriately to neighbors with the same first three components.
- $nfts$, a non-positive timestamp whose absolute value is the causal clock time when the current leader was elected.
- lid , the id of the current leader.

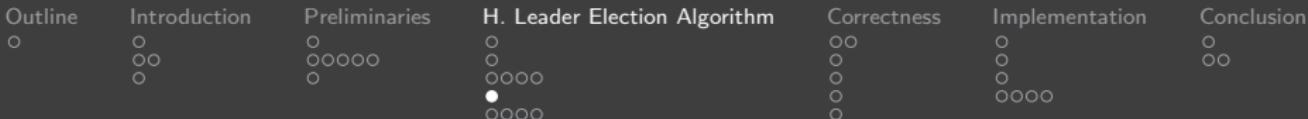




Nodes, Neighbors and Heights

- δ , an integer that is set to ensure that links are directed appropriately to neighbors with the same first three components.
- $nfts$, a non-positive timestamp whose absolute value is the causal clock time when the current leader was elected.
- lid , the id of the current leader.
- id , the node's unique ID.



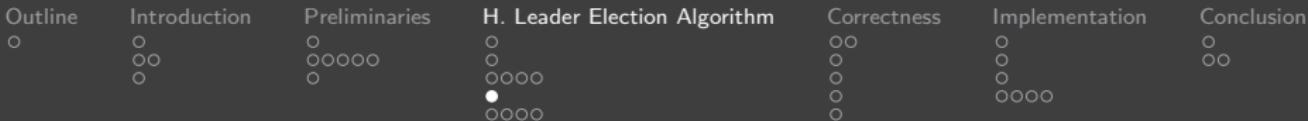


Initial State

Initial State

- $forming_u$ is empty,
- N_u equals the set of neighbors of u in G_{chan}^{init}



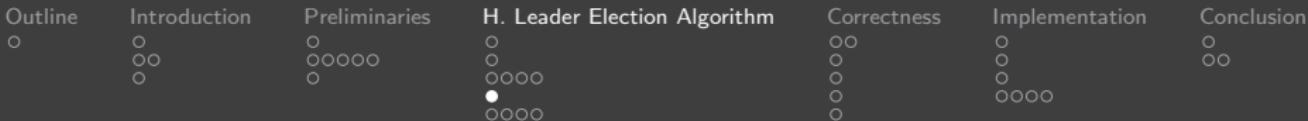


Initial State

Initial State

- $forming_u$ is empty,
- N_u equals the set of neighbors of u in G_{chan}^{init}
- $height_u[u] = (0, 0, 0, \delta_u, 0, l, u)$ where l is the *id* of a fixed node in u 's connected component in G_{chan}^{init} (the current leader), and δ_u equals the distance from u to l in G_{chan}^{init} ,



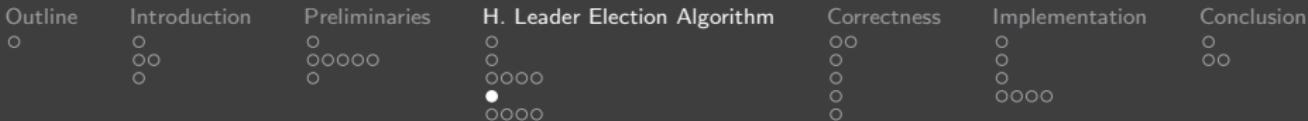


Initial State

Initial State

- $forming_u$ is empty,
- N_u equals the set of neighbors of u in G_{chan}^{init}
- $height_u[u] = (0, 0, 0, \delta_u, 0, l, u)$ where l is the *id* of a fixed node in u 's connected component in G_{chan}^{init} (the current leader), and δ_u equals the distance from u to l in G_{chan}^{init} ,
- for each v in N_u , $height_u[v] = height_v[v]$ (i.e., u has accurate information about v 's height), and



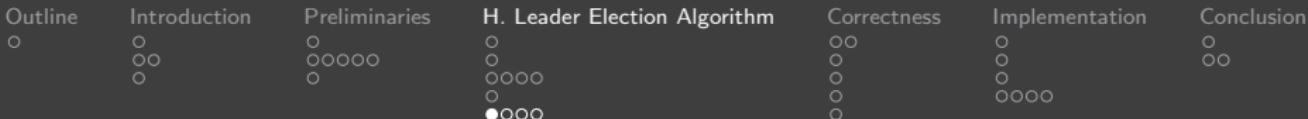


Initial State

Initial State

- $forming_u$ is empty,
- N_u equals the set of neighbors of u in G_{chan}^{init}
- $height_u[u] = (0, 0, 0, \delta_u, 0, l, u)$ where l is the *id* of a fixed node in u 's connected component in G_{chan}^{init} (the current leader), and δ_u equals the distance from u to l in G_{chan}^{init} ,
- for each v in N_u , $height_u[v] = height_v[v]$ (i.e., u has accurate information about v 's height), and
- \mathcal{T}_u is initialized properly with respect to the definition of causal clocks.





Algorithm Description

Algorithm 1 When $ChannelDown_{uv}$ event occurs:

```

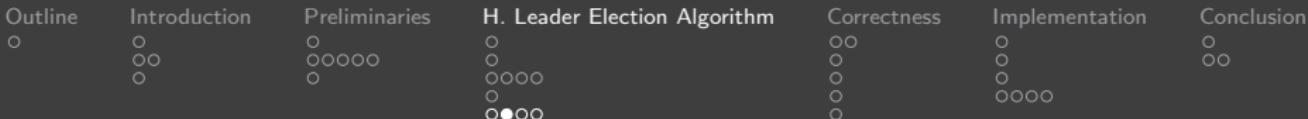
1:  $N := N \setminus v$ 
2:  $forming := forming \setminus v$ 
3: if  $N = \emptyset$  then
4:   ELECTSELF
5:   send Update(height[u]) to all  $w \in forming$ )
6: else if SINK then
7:   STARTNEWREFLEVEL
8:   send Update(height[u]) to all  $w \in (N \cup forming)$ )
9: else if  $id^v = pred^u$  and  $id^v = slid^u$  then
10:    $pred^u = \min \{k \mid k \in N^u \text{ and } \delta^k = \delta^u\}$ 
11:    $slid^u = \min \{k \mid k \in N^u \text{ and } \delta^k = \delta^u\}$ 
12: else if  $id^v = pred^u$  and  $id^v \neq slid^u$  then
13:    $pred^u = \min \{k \mid k \in N_u \text{ and } \delta^k = \delta^u\}$ 
14:    $slid^u = slid^{\min \{k \mid k \in N_u \text{ and } \delta^k = \delta^u\}}$ 
15: end if
  
```

Algorithm 2 When $ChannelUp_{uv}$ event occurs:

```

1:  $forming := forming \cup v$ 
2: send Update(height[u]) to  $v$ 
  
```

Figure: Code triggered by Topology Changes



Algorithm Description

Algorithm 3 When node u receives $Update(h)$ from node $v \in forming \cup N$:

```

1:  $height[v] := h$                                 ▷ if  $v$  is in neither forming nor  $N$ , message is ignored
2:  $forming := forming \setminus v$ 
3:  $N := N \cup v$ 
4:  $myOldHeight := height[u]$ 
5: if  $(nlts_u, lid_u) = (nlts_v, lid_v)$            ▷ leader pair are the same
6:   if  $SINK$  then
7:     if  $(\exists(\tau, oid, r) \mid (\tau_w, oid_w, r_w) = (\tau, oid, r) \forall w \in N)$  then
8:       if  $(\tau > 0)$  and  $(r = 0)$  then
9:         REFLECTREFLEVEL
10:        else if  $(\tau > 0)$  and  $(r = 1)$  and  $(oid = u)$  then
11:          ELECTSELF
12:        else                                     ▷  $(\tau = 0)$  or  $(\tau > 0$  and  $r = 1$  and  $oid \neq u)$ 
13:          STARTNEWREFLEVEL
14:        end if
15:      else                                     ▷ neighbors have different ref levels
16:        PROPAGATELARGESTREFLEVEL
17:      end if
18:    end if                                     ▷ else not sink, do nothing
19:  else                                     ▷ leader pairs are different
20:    ADOPTLPIFPRIORITY( $v$ )
21:  end if
22:  if  $myOldHeight \neq height[u]$  then
23:    send  $Update(height[u])$  to all  $w \in (N \cup forming)$ 
24:  end if

```



Figure: Code triggered by Update Message



Algorithm Description

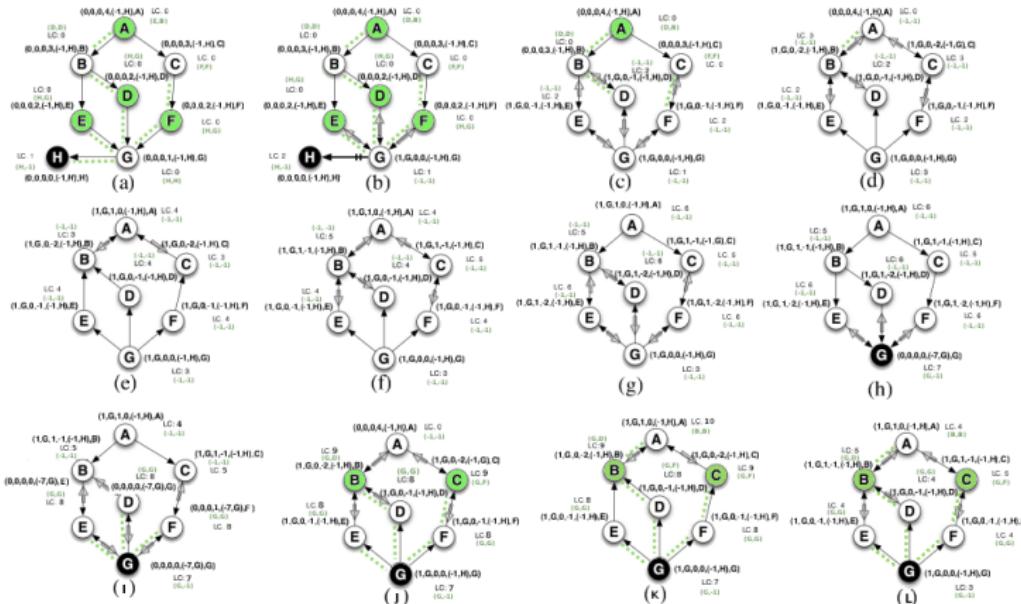
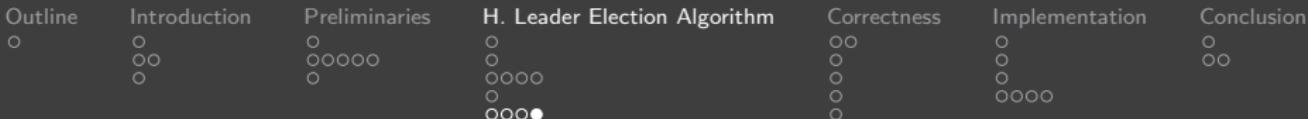


Figure: Sample Execution





Algorithm Description

Algorithm 4 ELECTSELF

1: $height[u] := (0, 0, 0, 0, -\mathcal{T}_u, u, u)$
 2: $SLP_u := (-1, -1)$

Algorithm 5 REFLECTREFLEVEL

1: $height[u] := (\tau, oid, 1, 0, nlts^u, lid^u, u)$
 2: $SLP_u := (-1, -1)$

Algorithm 6 PROPAGATELARGESTREFLEVEL

1: $(\tau^u, oid^u, r^u) := \max \{(\tau^k, oid^k, r^k) \mid k \in N\}$
 2: $\delta^u := \min \{\delta^k \mid k \in N \text{ and } (\tau^u, oid^u, r^u) = (\tau^k, oid^k, r^k)\} - 1$
 3: $SLP_u := (-1, -1)$

Algorithm 7 STARTNEWREFLEVEL

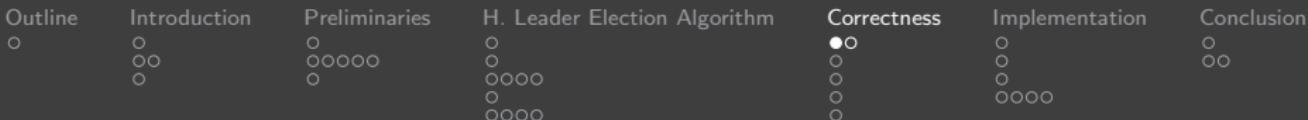
1: $height[u] := (\tau, oid, 1, 0, nlts^u, lid^u, u)$
 2: $SLP_u := (-1, -1)$

Algorithm 8 ADOPTLPIFPRIORITY(v)

1: if $nlts^v < nlts^u$ or $nlts^v = nlts^u$ and $lid^v < lid^u$ or $nlts^v = nlts^u$ and $lid^v = lid^u$ and $id^v < pred^u$ then
 2: $height[u] := (\tau^v, oid^v, r^v, \delta^v + 1, nlts^v, lid^v, u)$
 3: if $(\delta^v \bmod D \neq 0)$ then
 4: $SLP_u = (slid^v, id^v)$
 5: else
 6: $SLP_u = (id^v, id^v)$
 7: end if
 8: end if

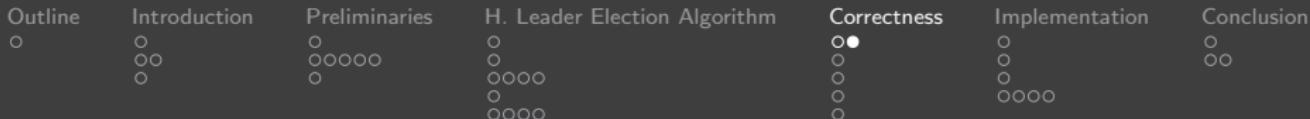


Figure: Subroutines



Outline

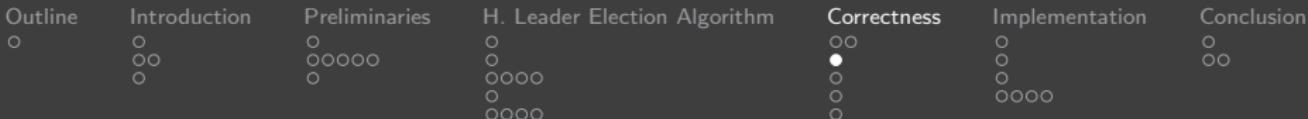
- 1** Introduction
 - Leader Election Problem
 - State of art
- 2** Preliminaries
 - System Model
 - Problem Definition
- 3** H. Leader Election Algorithm
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Algorithm Description
- 4** Correctness
 - Bounding the Number of Elections
 - Bounding the Number of New Reference Levels
 - Bounding the Number of Messages
 - Proving The Leader-Oriented DAG Aspect
- 5** Implementation
 - The Tool Used
 - Simulation
 - Performance Test
- 6** Conclusion
 - Is The Algorithm Perfect ?



How the proof is dealt with ?

The proof uses a number of invariants, denoted as “Properties”, which are shown to hold in every configuration of every execution; each one is proved (separately) by induction on the configurations occurring in an execution.



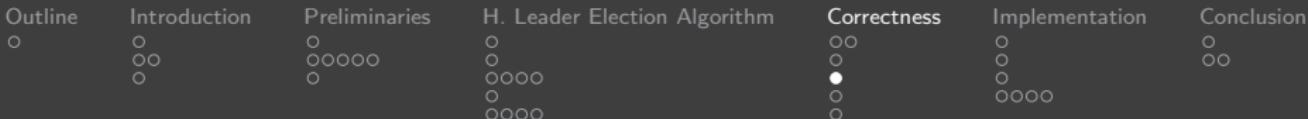


Bounding the Number of Elections

Bounding the Number of Elections

We bound, in Lemma 3, the number of elections that can occur after the last topology change; this result relies on the fact, shown in Lemma 2, that once a node u adopts a leader that was elected after the last topology change, u never becomes a sink again.



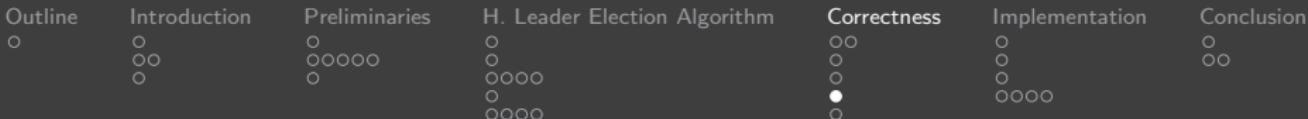


Bounding the Number of New Reference Levels

Bounding the Number of Elections

We bound, in Lemma 4, the number of new reference levels that are started after the last topology change; the proof of this result relies on several additional properties.



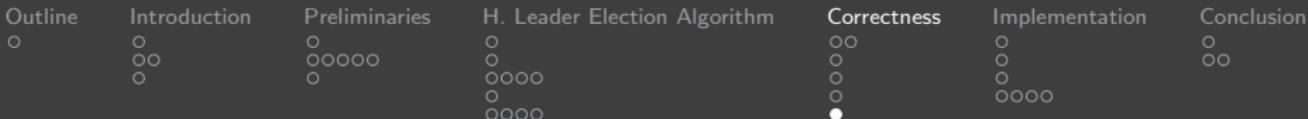


Bounding the Number of Messages

Bounding the Number of Messages

We show, in Lemmas 5, 6, and 7, that eventually there are no messages in transit and every node has an accurate view of its neighbors' heights.



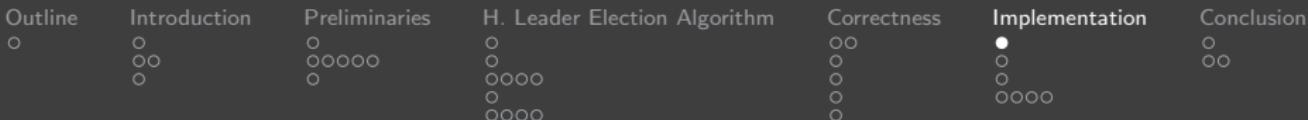


Proving The Leader-Oriented DAG Aspect

Proving The Leader-Oriented DAG Aspect

All the pieces are put together in Theorem 1 of Section 4.6 to show that eventually we have a leader-oriented connected component; a couple of additional properties are needed for this result.





Outline

1 Introduction

- Leader Election Problem
- State of art

2 Preliminaries

- System Model
- Problem Definition

3 H. Leader Election Algorithm

- Informal Description
- Nodes, Neighbors and Heights
- Initial State
- Algorithm Description

4 Correctness

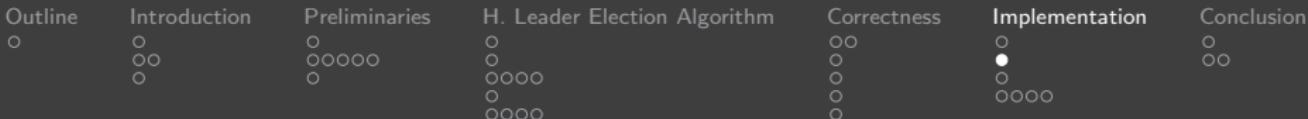
- Bounding the Number of Elections
- Bounding the Number of New Reference Levels
- Bounding the Number of Messages
- Proving The Leader-Oriented DAG Aspect

5 Implementation

- The Tool Used
- Simulation
- Performance Test

6 Conclusion

- Is The Algorithm Perfect ?

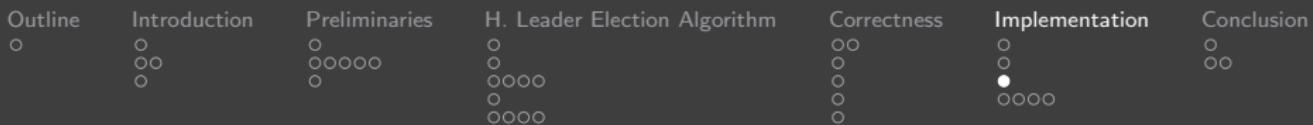


The Tool Used

What's JBotSim ?

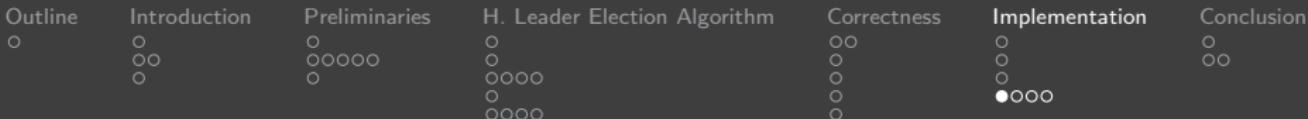
JBoTSim is a java library that offers basic primitives for proto-typing, running, and visualizing distributed algorithms in dynamic networks.





Simulation



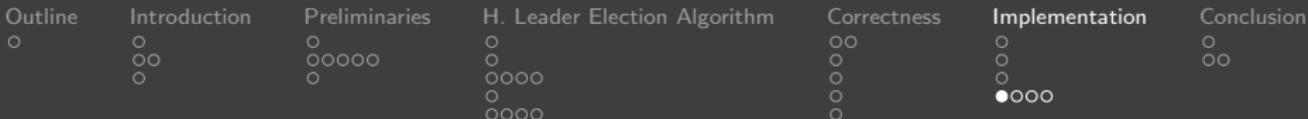


Performance Test

Performance Characteristics

- 1 Latency** / The number of rounds necessary for a component to elect a leader (become stable) after a topology change (average over all initial topologies).



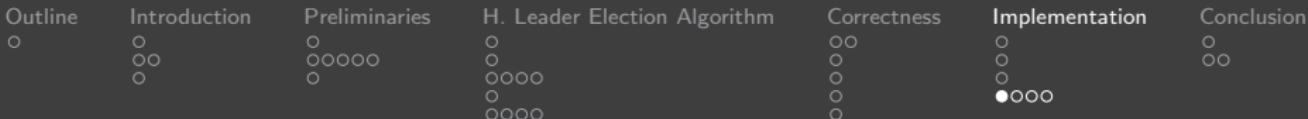


Performance Test

Performance Characteristics

- 1 Latency** l . The number of rounds necessary for a component to elect a leader (become stable) after a topology change (average over all initial topologies).
- 2 Sensitivity** s . The number of nodes that updated their heights in response to a topology change (average over all possible topology changes).



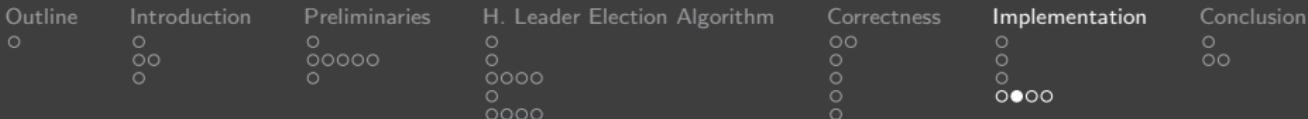


Performance Test

Performance Characteristics

- 1 Latency** ℓ . The number of rounds necessary for a component to elect a leader (become stable) after a topology change (average over all initial topologies).
- 2 Sensitivity** s . The number of nodes that updated their heights in response to a topology change (average over all possible topology changes).
- 3 Resilience** r . The maximal fraction of links which can go down in a stable component without reelecting the current leader.



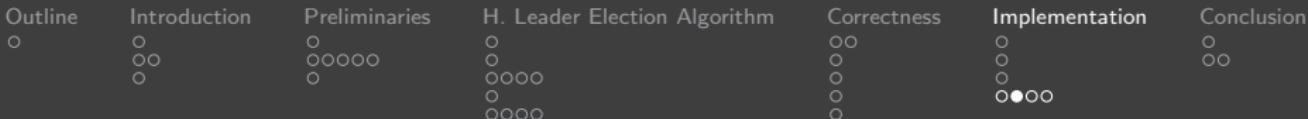


Performance Test

Performance Characteristics

- Merging two stable components:



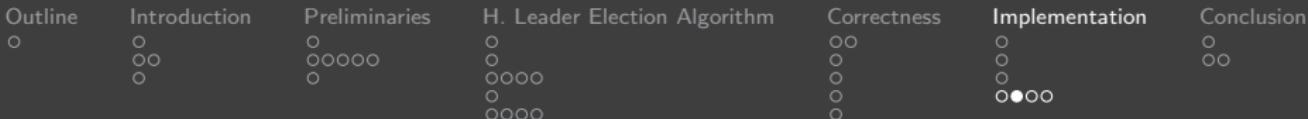


Performance Test

Performance Characteristics

- Merging two stable components:
 - Given two fully connected (every two nodes are neighbours) components: $l \simeq 2$.



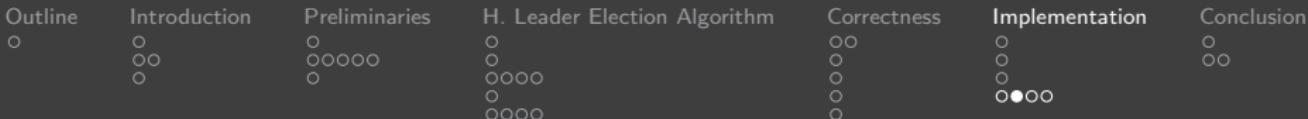


Performance Test

Performance Characteristics

- Merging two stable components:
 - Given two fully connected (every two nodes are neighbours) components: $l \simeq 2$.
 - Given two poorly connected (a node can have not more than two neighbours) components: $l \simeq n$.



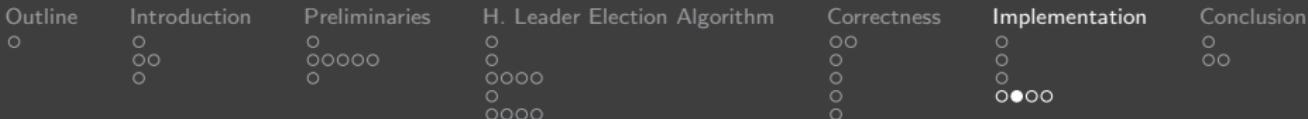


Performance Test

Performance Characteristics

- Merging two stable components:
 - Given two fully connected (every two nodes are neighbours) components: $l \simeq 2$.
 - Given two poorly connected (a node can have not more than two neighbours) components: $l \simeq n$.
- Partitioning a stable component:



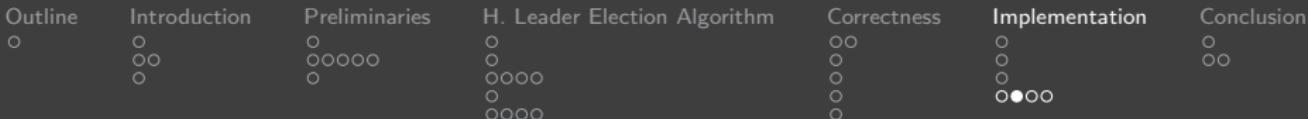


Performance Test

Performance Characteristics

- Merging two stable components:
 - Given two fully connected (every two nodes are neighbours) components: $l \simeq 2$.
 - Given two poorly connected (a node can have not more than two neighbours) components: $l \simeq n$.
- Partitioning a stable component:
 - Getting two fully connected component: $l = 2$.



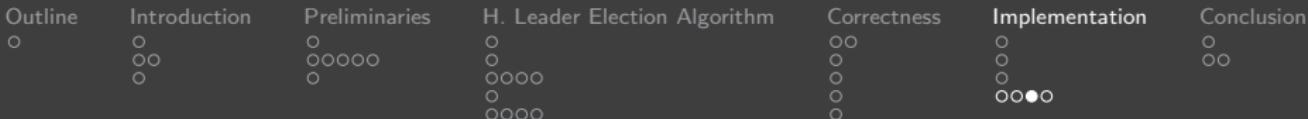


Performance Test

Performance Characteristics

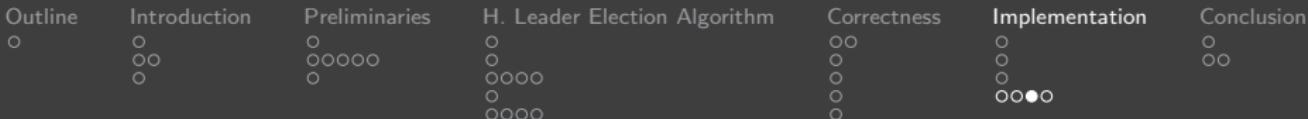
- Merging two stable components:
 - Given two fully connected (every two nodes are neighbours) components: $l \simeq 2$.
 - Given two poorly connected (a node can have not more than two neighbours) components: $l \simeq n$.
- Partitioning a stable component:
 - Getting two fully connected component: $l = 2$.
 - Getting two poorly connected component: $l = 2n$.





- Topology change in a “small world-like” stable component (every node has $O(\log n)$ neighbours) without partitioning:

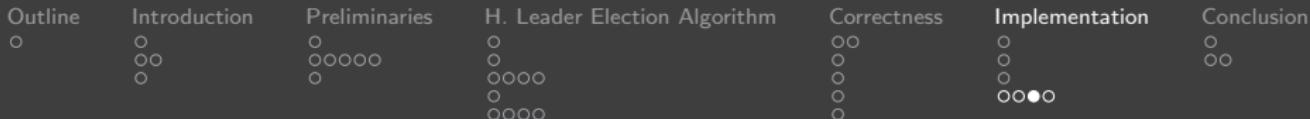




Performance Test

- Topology change in a “small world-like” stable component (every node has $O(\log n)$ neighbours) without partitioning:
 - Latency $l = O(1)$.

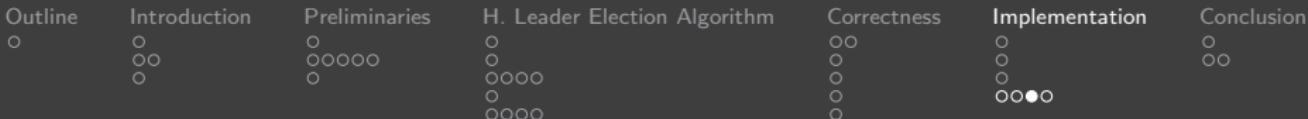




Performance Test

- Topology change in a “small world-like” stable component (every node has $O(\log n)$ neighbours) without partitioning:
 - Latency $l = O(1)$.
 - Sensitivity $s = O(\log(n))$.

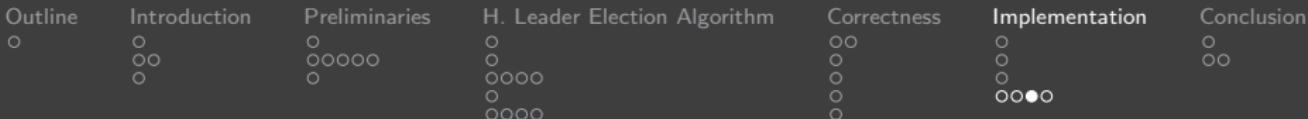




Performance Test

- Topology change in a “small world-like” stable component (every node has $O(\log n)$ neighbours) without partitioning:
 - Latency $l = O(1)$.
 - Sensitivity $s = O(\log(n))$.
- Resilience of a stable component:

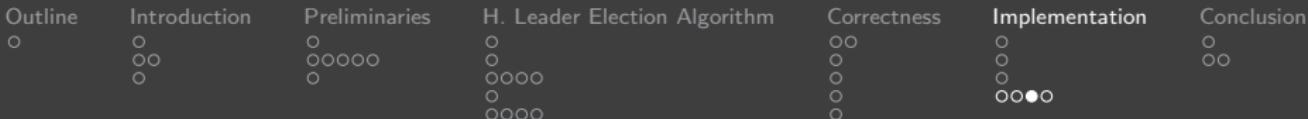




Performance Test

- Topology change in a “small world-like” stable component (every node has $O(\log n)$ neighbours) without partitioning:
 - Latency $l = O(1)$.
 - Sensitivity $s = O(\log(n))$.
- Resilience of a stable component:
 - Given a fully connected component: $r \simeq 1 - 2/n$.

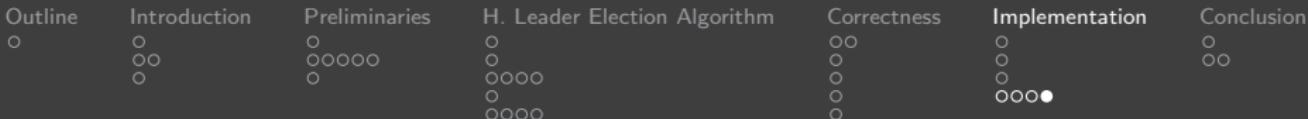




Performance Test

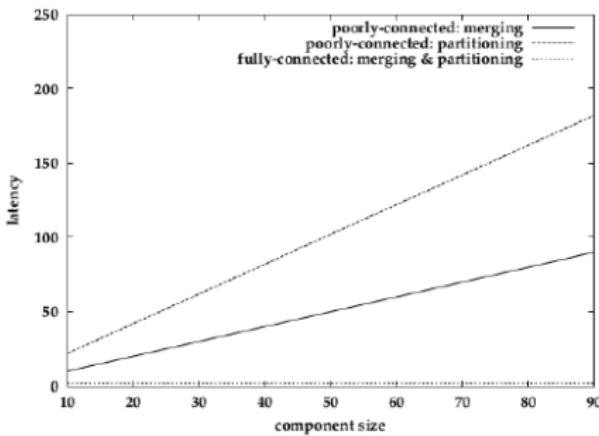
- Topology change in a “small world-like” stable component (every node has $O(\log n)$ neighbours) without partitioning:
 - Latency $l = O(1)$.
 - Sensitivity $s = O(\log(n))$.
- Resilience of a stable component:
 - Given a fully connected component: $r \simeq 1 - 2/n$.
 - Given a “small world-like” component: $r \simeq 1 - 2/\log(n)$.



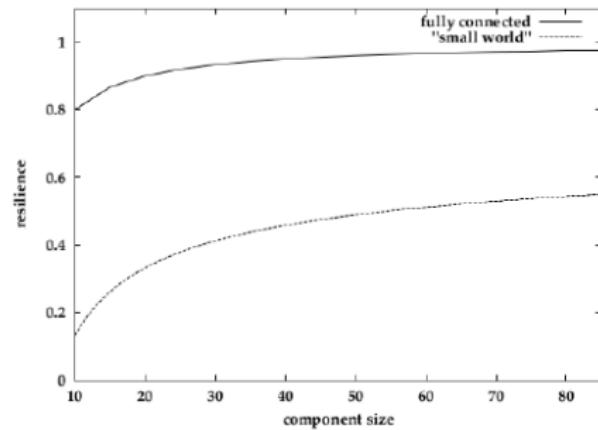


Performance Test

How is our algorithm's performance ?



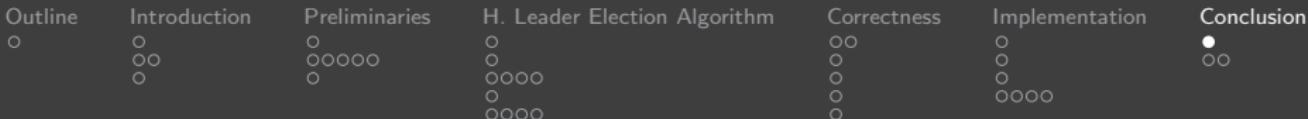
(a) Latency



(b) Resilience

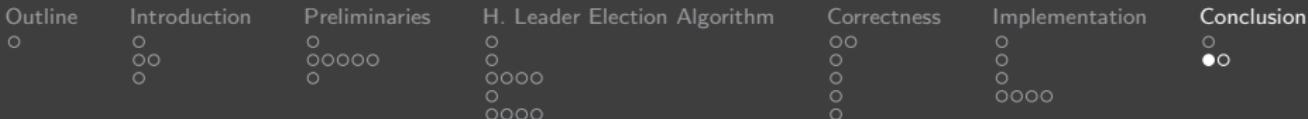


Figure: Simulation Results



Outline

- 1** Introduction
 - Leader Election Problem
 - State of art
- 2** Preliminaries
 - System Model
 - Problem Definition
- 3** H. Leader Election Algorithm
 - Informal Description
 - Nodes, Neighbors and Heights
 - Initial State
 - Algorithm Description
- 4** Correctness
 - Bounding the Number of Elections
 - Bounding the Number of New Reference Levels
 - Bounding the Number of Messages
 - Proving The Leader-Oriented DAG Aspect
- 5** Implementation
 - The Tool Used
 - Simulation
 - Performance Test
- 6** Conclusion
 - Is The Algorithm Perfect ?

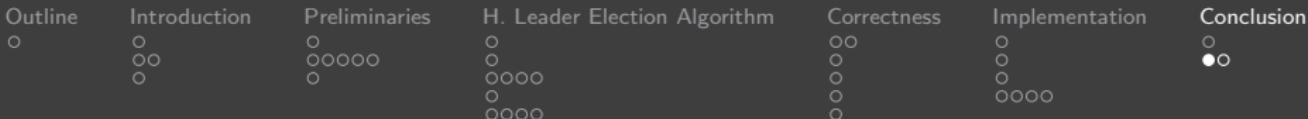


Is The Algorithm Perfect ?

Is The Algorithm Perfect ?

- An open question is how to extend our algorithm and its analysis to handle a wider range of clocks, such as approximately synchronized clocks and vector clocks.



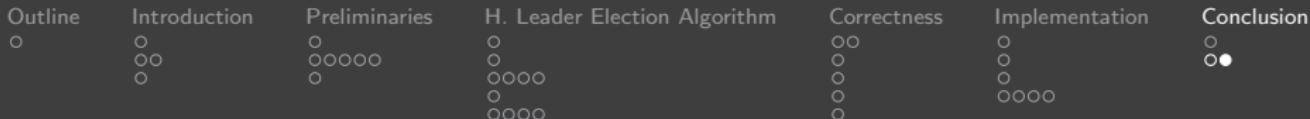


Is The Algorithm Perfect ?

Is The Algorithm Perfect ?

- An open question is how to extend our algorithm and its analysis to handle a wider range of clocks, such as approximately synchronized clocks and vector clocks.
- Another question could be asked concerning the possibility of building up an optimized spanning tree using some algorithms such as Kruskal's Minimum Spanning Tree Algorithm (MST).





Is The Algorithm Perfect ?

Thank you for your attention

All your questions and remarks are welcomed.

